

Parallel Programming: Moore's Law and Multicore



Oregon State
University
Mike Bailey

mjb@cs.oregonstate.edu



This work is licensed under a [Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License](https://creativecommons.org/licenses/by-nc-nd/4.0/)

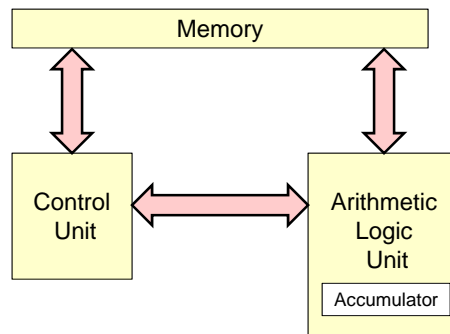


Oregon State
University
Computer Graphics

moores.law.and.multicore.pptx

mjb - March 15, 2020

Von Neumann Architecture: Basically the fundamental pieces of a CPU have not changed since the 1960s



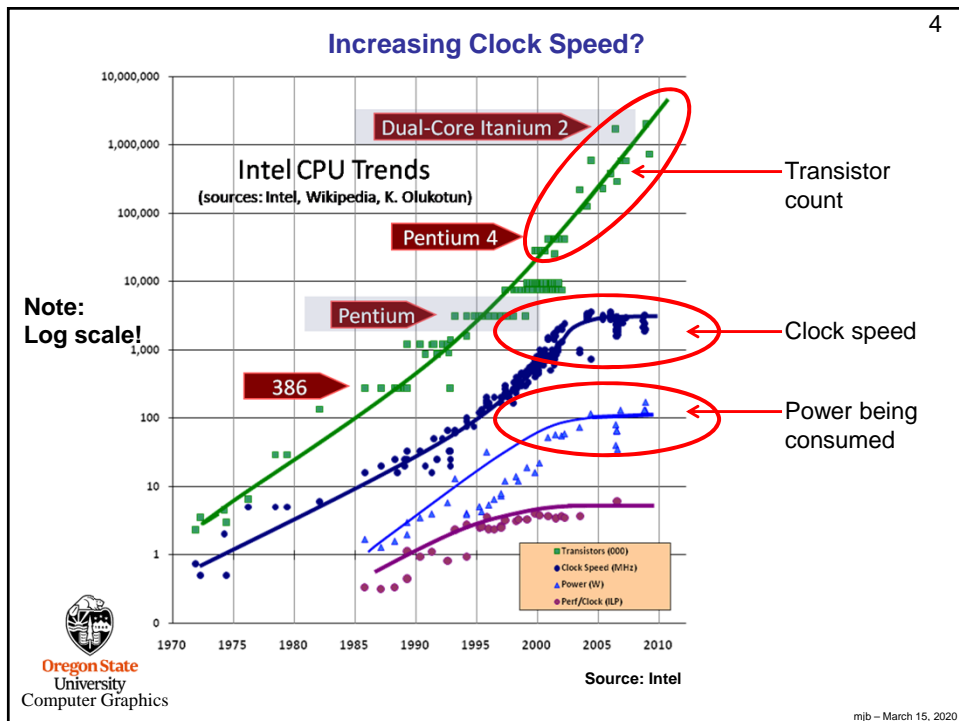
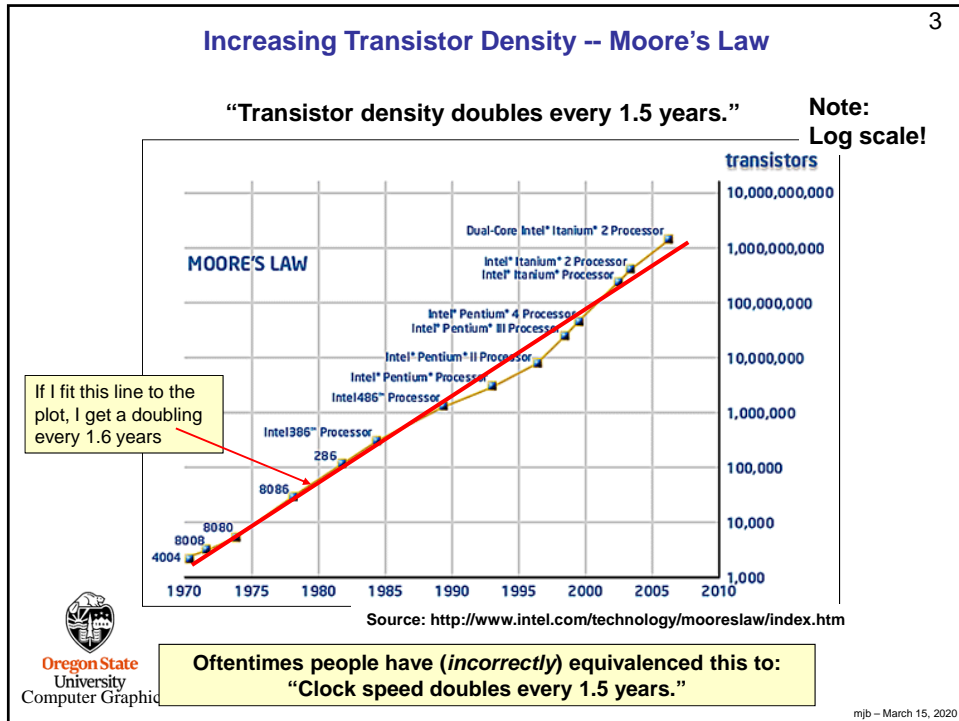
Other elements:

- Clock
- Registers
- Program counter
- Stack pointer



Oregon State
University
Computer Graphics

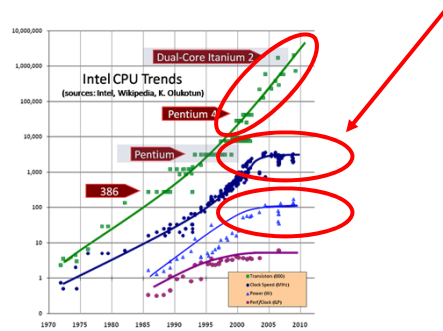
mjb - March 15, 2020



Moore's Law

5

- Fabrication process size ("gate pitch") has fallen from 65 nm, to 45 nm, to 32 nm, to 22 nm, to 16 nm, to 11 nm, to 8 nm. This translates to more transistors on the same size die.
- From 1986 to 2002, processor performance increased an average of 52%/year, but then virtually plateaued.



Clock Speed and Power Consumption

6

1981	IBM PC	5 MHz
1995	Pentium	100 MHz
2002	Pentium 4	3000 MHz (3 GHz)
2007		3800 MHz (3.8 GHz)
2009		4000 MHz (4.0 GHz)

Clock speed has hit a plateau, largely because of power consumption and dissipation.

$$\text{PowerConsumption} \propto \text{ClockSpeed}^2$$

is-proportional-to

Yikes!

Once consumed, that power becomes *heat*, which much be dissipated somehow. In general, compute systems can remove around 150 watts/cm without resorting to exotic cooling methods.

And, speaking of “exotic”, AMD set the world record for clock speed (8.429 GHz) using a Liquid Nitrogen-cooled CPU

7



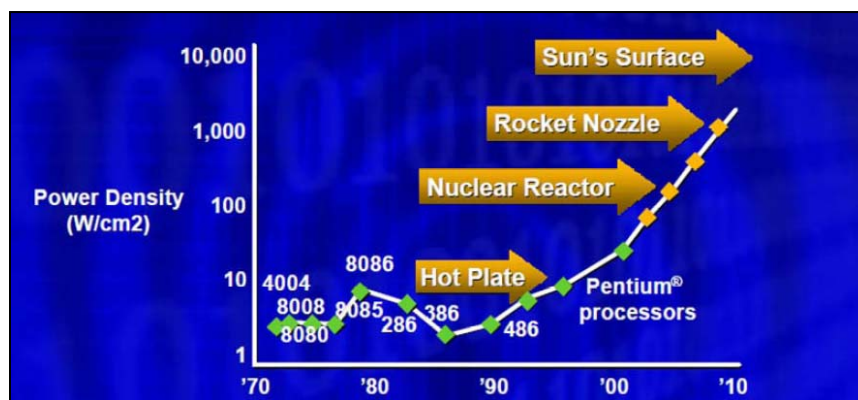
Oregon State
University
Computer Graphics

Source: AMD

mjb - March 15, 2020

What Kind of Power Density Dissipation Would it Have Taken to Keep up with Clock Speed Trends?

8



Source: Intel

Oregon State
University
Computer Graphics

mjb - March 15, 2020

MultiCore -- Multiprocessing on a Single Chip

9

So, to summarize:

Moore's Law of transistor density is still going, but the "Moore's Law" of clock speed has hit a wall. Now what do we do?

We keep packing more and more transistors on a single chip, but don't increase the clock speed. Instead, we increase computational throughput by using those transistors to pack multiple processors onto the same chip.

This is referred to as **multicore**.



Oregon State
University
Computer Graphics

Vendors have also reacted by adding SIMD floating-point units on the chip as well. We will get to that later.

mjb - March 15, 2020


MultiCore and Multithreading

10

Multicore, even without multithreading too, is still a good thing. It can be used, for example, to allow multiple programs on a desktop system to always be executing concurrently.

Multithreading, even without multicore too, is still a good thing. Threads can make it easier to logically have many things going on in your program at a time, and can absorb the dead-time of other threads.

But, the big gain in performance is to use *both* to speed up a *single program*. For this, we need a **combination of both multicore and multithreading**.



Multicore is a very hot topic these days. It would be hard to buy a CPU that doesn't have more than one core. We, as programmers, get to take advantage of that.




Oreg
University
Computer Graphics

We need to be prepared to convert our programs to run on **MultiThreaded Shared Memory Multicore** architectures.

mjb - March 15, 2020

Each of the Multiple Cores keeps its own State

11

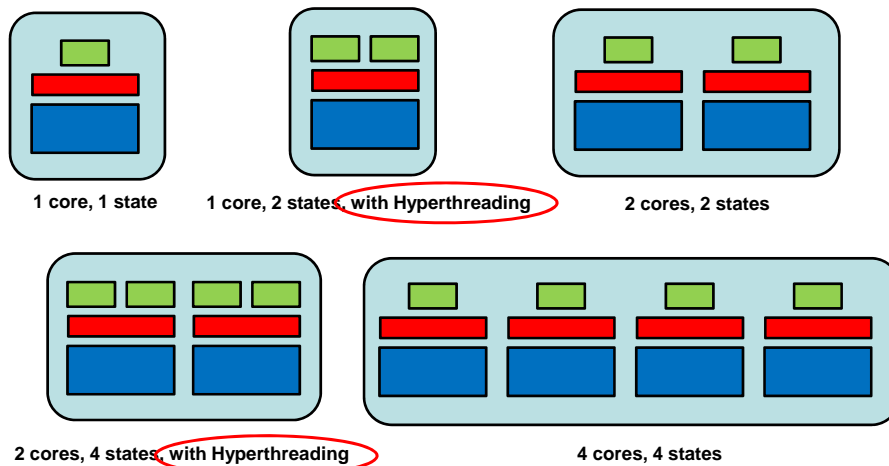


- Registers
- Program Counter
- Stack Pointer

mjb - March 15, 2020

So, if that's what Multicore is about, what is *Hyperthreading*?

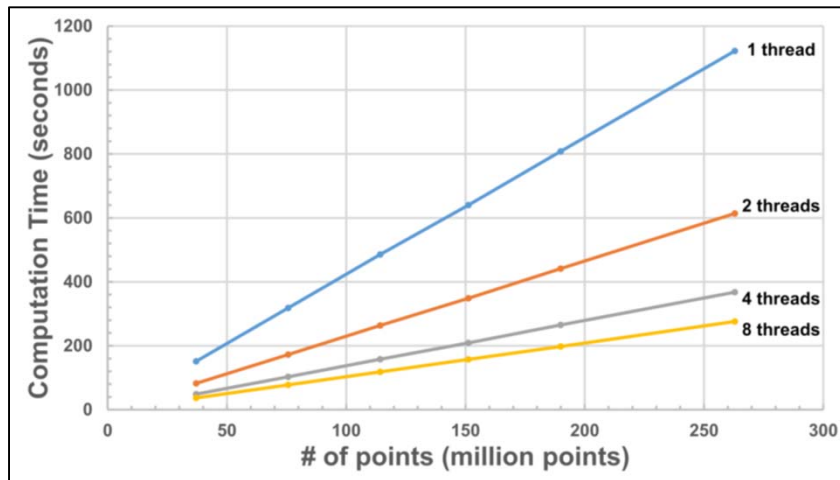
12



mjb - March 15, 2020

Four Cores with Two Hyperthreads per Core

13



Source: Erzhua Che



Note that this is upside-down from our usual convention. Sorry. I got this from someone else.



mjb - March 15, 2020