

OpenMP Tasks



Oregon State
University
Mike Bailey


mjb@cs.oregonstate.edu



This work is licensed under a [Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License](https://creativecommons.org/licenses/by-nc-nd/4.0/)



Oregon State
University
Computer Graphics



tasks.pptx

mjb - March 23, 2020

Remember OpenMP Sections?

Sections are independent blocks of code, able to be assigned to separate threads if they are available.

```
#pragma omp parallel sections
{
    #pragma omp section
    {
        Task 1
    }
    #pragma omp section
    {
        Task 2
    }
}
```

There is an **implied barrier** at the end



Oregon State
University
Computer Graphics

OpenMP sections are **static**, that is, they are good if you know, when you are writing the program, how many of them you will need.

mjb - March 23, 2020

It would be nice to have something more Dynamic

3



Imagine a capability where you can write something to do down on a Post-It® note, accumulate the Post-It notes, then have all of the threads together execute that set of tasks.

You would also like to not have to know, ahead of time, how many of these Post-It notes you will write. That is, you want the total number to be **dynamic**.

Well, congratulations, you have just invented **OpenMP Tasks**!

Oregon State University
Computer Graphics

mjb - March 23, 2020

OpenMP Tasks

4

- An OpenMP task is a single line of code or a structured block which is immediately “written down” in a list of tasks.
- The new task can be executed immediately, or it can be deferred.
- If the *if* clause is used and the argument evaluates to 0, then the task is executed immediately, superseding whatever else that thread is doing.
- There has to be an existing parallel thread team for this to work. Otherwise one thread ends up doing all tasks and you don’t get any contribution to parallelism.
- One of the best uses of this is to process elements of a linked list or a tree.

You can create a task barrier with:

#pragma omp taskwait

Tasks are very much like OpenMP **Sections**, but Sections are static, that is, the number of sections is set when you write the code, whereas **Tasks** can be created anytime, and in any number, under control of your program’s logic.

Oregon State University
Computer Graphics

mjb - March 23, 2020

OpenMP Task Example: Something (Supposedly) Simple

5


```
omp_set_num_threads( 2 );
#pragma omp parallel default(none)
{
    #pragma omp task
    fprintf( stderr, "A\n" );

    #pragma omp task
    fprintf( stderr, "B\n" );
}
```

Without this, thread #0 has to do everything

Writes fprintf(stderr, "A\n"); on a sticky note and adds it to the list

Writes fprintf(stderr, "B\n"); on a sticky note and adds it to the list



If You Run This a Number of Times, You Get This: (What Happened?)

6

Run #	1	2	3	4	5
	B	B	B	B	B
	A	B	A	A	A
	B	A	A	A	B
	A	A	B	B	A

1. Why do we not get the same output every time?
2. Why do we get 4 things printed when we only have print statements in 2 tasks?

Not so simple, huh?

The first answer is easy. Unless you make some special arrangements, the order of execution of the different tasks is *undefined*.

The second answer is that we actually asked each of the two threads to put two tasks on the sticky notes, for a total of four. How can we get only one thread to do this?

The “single” Pragma

7


```
omp_set_num_threads( 2 );
#pragma omp parallel default(none)
{
    #pragma omp single
    {
        #pragma omp task
        fprintf( stderr, "A\n" );

        #pragma omp task
        fprintf( stderr, "B\n" );
    }
}
```

When using Tasks, you only want *one* thread to write the things to do down on the sticky note, but you want *all* of the threads to be able to execute the sticky notes.



Oregon State
University
Computer Graphics



mjb - March 23, 2020

But, if you run this, the order of printing will still be non-deterministic. To solve that problem, do this:

8

```
omp_set_num_threads( 2 );
#pragma omp parallel
{
    #pragma omp single default(none)
    {
        #pragma omp task
        fprintf( stderr, "A\n" );


        #pragma omp taskwait ← Causes all tasks to wait until they are completed

        #pragma omp task
        fprintf( stderr, "B\n" );

        #pragma omp taskwait ← Causes all tasks to wait until they are completed
    }
}
```



Oregon State
University
Computer Graphics



mjb - March 23, 2020

A Better OpenMP Task Example: Processing each Element of a Linked List

9

Without this, thread #0 has to do everything

```
#pragma omp parallel default(none)
{
    Without this, each thread does a full traversal – bad idea!
    #pragma omp single default(none)
    {
        element *p = listHead;
        while( p != NULL )
        {
            Write "Process( p )" on a sticky note and add it to the list
            #pragma omp task firstprivate(p)
            Process( p );
            p = p->next;
        }
    }
    #pragma omp taskwait
}
```

Put this here if you want to wait for all tasks to finish being executed before proceeding

University
Computer Graphics

mjb – March 23, 2020

One more thing – Task Dependencies

10

Remember from before: unless you make some special arrangements, the order of execution of the different tasks is *undefined*. Here come the special arrangements.

```
omp_set_num_threads( 3 );
#pragma omp parallel
{
    #pragma omp single default(none)
    {
        float a, b, c;
        #pragma omp task depend( OUT: a )
        a = 10.;

        #pragma omp task depend( IN: a, OUT: b )
        b = a + 16.;

        #pragma omp task depend( IN: b )
        c = b + 12.;
    }
    #pragma omp taskwait
}
```



This maintains the proper dependencies, but, because it involves all of the tasks, it essentially serializes the parallelism out of them.
Be careful not to go overboard with dependencies!




mjb – March 23, 2020

Tree Traversal Algorithms

11

Given a tree:



- We would like to traverse it as quickly as possible.
- We are assuming that we do not need to traverse it in order.
- We just need to visit all nodes.

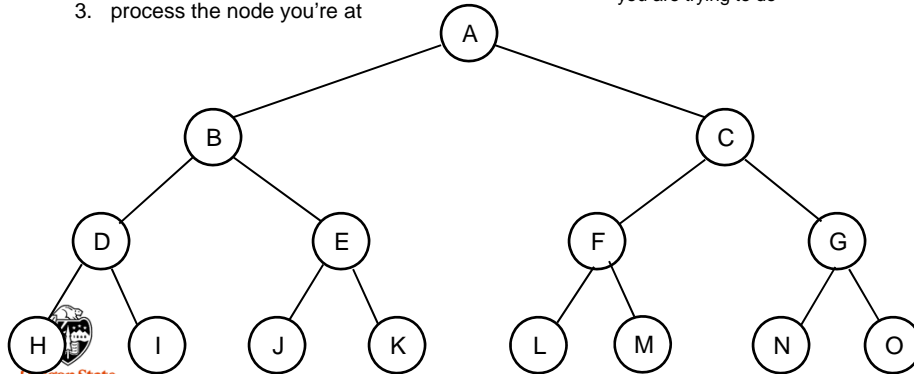
mjb - March 23, 2020

Tree Traversal Algorithms

12

- This is common in graph algorithms, such as searching.
- If the tree is binary and is balanced, then the maximum depth of the tree is $\log_2(\# \text{ of Nodes})$
- Strategy at a node:
 1. follow one descendent node
 2. follow the other descendent node
 3. process the node you're at

This order could be re-arranged, depending on what you are trying to do



mjb - March 23, 2020

Tree Traversal Algorithms

13



```
#pragma omp parallel
```

```
#pragma omp single
```

```
Traverse( root );
```

```
#pragma omp taskwait
```

Without this, thread #0 has to do everything

Without this, each thread does a full traversal – bad idea!

Put this here if you want to wait for all nodes to be traversed before proceeding



mjb - March 23, 2020

Parallelizing a Binary Tree Traversal with Tasks

14



```
void
Traverse( Node *n )
{
    if( n->left != NULL )
    {
        #pragma omp task private(n) untied
        Traverse( n->left );
    }

    if( n->right != NULL )
    {
        #pragma omp task private(n) untied
        Traverse( n->right );
    }

    #pragma omp taskwait

    Process( n );
}
```

Put this here if you want to wait for both branches to be taken before processing the parent




mjb - March 23, 2020

Parallelizing a Binary Tree Traversal with Tasks

15

Traverse(A);

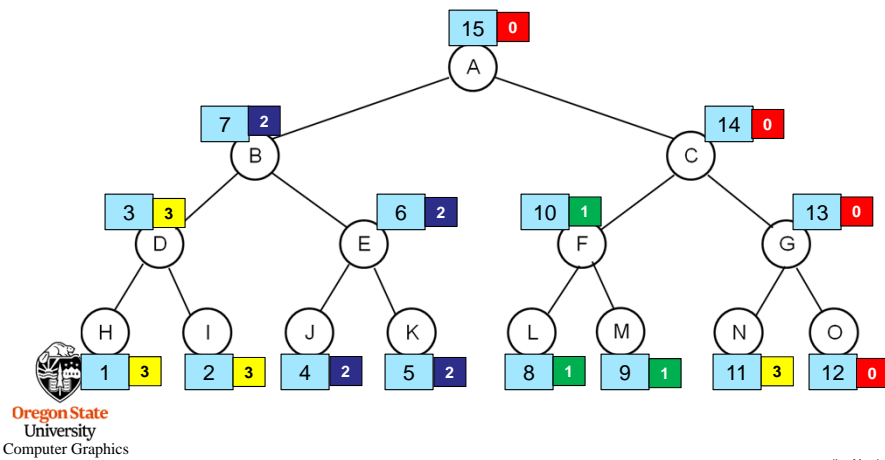


Parallelizing a Binary Tree Traversal with Tasks: *Tied* (gcc 8.2)

16

Threads:
0 1 2 3

Traverse(A);

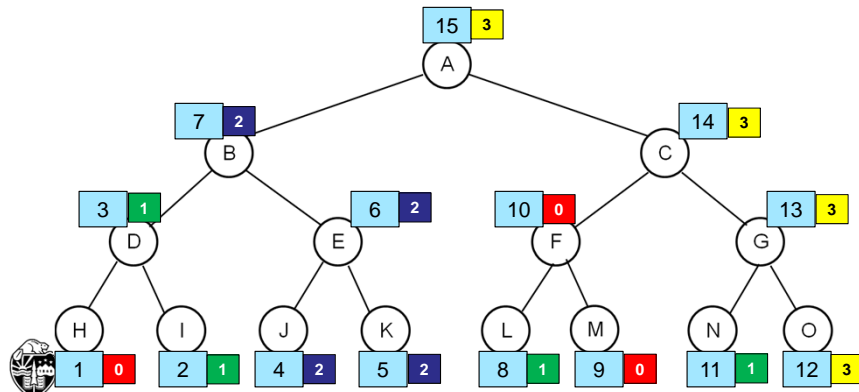


Parallelizing a Binary Tree Traversal with Tasks: *Untied* (gcc 8.2)

17

Threads:
0 1 2 3

Traverse(A);



Oregon State
University
Computer Graphics

mjb - March 23, 2020

How Evenly Tasks Get Assigned to Threads

18

6 Levels – g++ 4.9:

Thread #	Number of Tasks
0	1
1	32
2	47
3	47

6 Levels – g++ 8.2:

Thread #	Number of Tasks
0	31
1	46
2	32
3	18

12 Levels – g++ 4.9:

Thread #	Number of Tasks
0	2561
1	2
2	2813
3	2815

12 Levels – g++ 8.2:

Thread #	Number of Tasks
0	1
1	2048
2	3071
3	3071

Oregon State
University
Computer Graphics

mjb - March 23, 2020

How Evenly Tasks Get Assigned to Threads

19

6 Levels – g++ 8.2:

Thread #	Number of Tasks
0	31
1	46
2	32
3	18

6 Levels – icpc 15.0.0:

Thread #	Number of Tasks
0	29
1	31
2	41
3	26

12 Levels – g++ 8.2:

Thread #	Number of Tasks
0	1
1	2048
2	3071
3	3071

12 Levels – icpc 15.0.0:

Thread #	Number of Tasks
0	1999
1	2068
2	2035
3	2089



Oregon State
University
Computer Graphics

mjb – March 23, 2020

Benchmarking a Binary Task-driven Tree Traversal

20

```
void
Process( Node *n )
{
    for( int i = 0; i < 1024; i++ )
    {
        n->value = pow( n->value, 1.1 );
    }
}
```




Oregon State
University
Computer Graphics

mjb – March 23, 2020


Performance vs. Number of Threads

21



Performance vs. Number of Levels

22



Parallelizing a Tree Traversal with Tasks

23

- Tasks get spread among the current “thread team”
- Tasks can execute immediately or can be deferred. They are executed at “some time”.
- Tasks can be moved between threads, that is, if one thread has a backlog of tasks to do, an idle thread can come steal some workload.
- Tasks are more dynamic than sections. The task paradigm would still work if there was a variable number of children at each node.

Parallelizing an N-Tree Traversal with Tasks

24

```
void
Traverse( Node *n )
{
    for( int i = 0; i < n->numChildren; i++ )
    {
        if( n->child[ i ] != NULL )
        {
            #pragma omp task
            Traverse( n->child[ i ] );
        }
    }

    #pragma omp taskwait

    Process( n );
}
```

