

The OSU College of Engineering DGX System for Advanced GPU Computing



**Oregon State
University**

Mike Bailey

mjb@cs.oregonstate.edu



This work is licensed under a [Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License](https://creativecommons.org/licenses/by-nc-nd/4.0/)



**Oregon State
University**
Computer Graphics

dgx_system.pptx

mjb – April 18, 2020

OSU's College of Engineering bought six Nvidia DGX-2 systems

Each DGX server:

- Has 16 NVidia Tesla V100 GPUs
- Has 28TB of disk, all SSD
- Has two 24-core Intel Xeon 8168 Platinum 2.7GHz CPUs
- Has 1.5TB of DDR4-2666 System Memory
- Runs the CentOS 7 Linux operating system

Overall compute power:

- Each V100 NVidia Tesla card has 5,120 CUDA Cores and 640 Tensor Cores
- This gives each 16-V100 DGX server a total of 81,920 CUDA cores and 10,240 Tensor cores
- This gives the entire 6-DGX package a total of 491,520 CUDA Cores and 61,440 Tensor Cores




**Oregon State
University**
Computer Graphics

mjb – April 18, 2020

Performance Comparison with one of our previous Systems

3



BTW, you can also use the *rabbit* machine:

ssh rabbit.engr.oregonstate.edu

It is a good place to write your code and get it to compile.

It is not a good place to run your code.

mjb - April 18, 2020

How to SSH to the DGX Systems

4

ssh over to a DGX submission machine --
submit-a and **submit-b** will also work

flip3 151% ssh submit-c.hpc.engr.oregonstate.edu

submit-c 142% module load slurm

Type this right away to set your path correctly



mjb - April 18, 2020

How to Check on the DGX Systems

5

submit-c 143% squeue

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
3923	mime4	c_only	jayasurw	R	1-10:32:19	1	compute-e-1
3963	mime4	2Dex	jayasurw	R	16:21:03	1	compute-e-2
3876	share	CH3COOH_	chukwuk	R	1-23:36:45	1	compute-2-6
3971	nerhp	tcsh	dionnec	R	8:59:45	1	compute-h-8
3881	dgx2	bash	heli	R	1-22:50:44	1	compute-dgx2-1
3965	dgx2	bash	chenju3	R	13:47:36	1	compute-dgx2-4
3645	dgx2	bash	mishrash	R	5-16:48:09	1	compute-dgx2-5
3585	dgx2	bash	azieren	R	6-17:34:00	1	compute-dgx2-3
3583	dgx2	bash	azieren	R	6-18:26:44	1	compute-dgx2-3

Check on the queues

submit-c 144% sinfo

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
share*	up	7-00:00:00	2	drain	compute-4-[3-4]
share*	up	7-00:00:00	1	mix	compute-2-6
sharegpu	up	7-00:00:00	1	mix	compute-dgxs-1
sharegpu	up	7-00:00:00	3	idle	compute-dgxs-[2-3],compute-gpu
dgx2	up	7-00:00:00	1	drain	compute-dgx2-2
dgx2	up	7-00:00:00	5	mix	compute-dgx2-[1,3-6]
gpu	up	7-00:00:00	2	mix	compute-gpu[3-4]
gpu	up	7-00:00:00	1	idle	compute-gpu2
gpu	up	7-00:00:00	1	down	compute-gpu1
dgx	up	7-00:00:00	3	mix	compute-dgx2-[4-6]
dgxs	up	7-00:00:00	1	mix	compute-dgxs-1
dgxs	up	7-00:00:00	2	idle	compute-dgxs-[2-3]
class	up	1:00:00	1	mix	compute-dgxs-1
class	up	1:00:00	2	idle	compute-dgxs-[2-3]
ecsc	up	7-00:00:00	1	mix	compute-2-6

System Information

Your partitions

Oregon State University
Computer Graphics

mjb - April 18, 2020

Submitting a CUDA job to the DGX Systems using Slurm

6

Create a shell file

submit.bash:

```
#!/bin/bash
#SBATCH -J MatrixMult
#SBATCH -A cs475-575
#SBATCH -p class
#SBATCH --gres=gpu:1
#SBATCH -o matrixmul.out
#SBATCH -e matrixmul.err
#SBATCH --mail-type=BEGIN,END,FAIL
#SBATCH --mail-user=joeparallel@oregonstate.edu
{ /usr/local/apps/cuda/cuda-10.1/bin/nvcc -o matrixMul matrixMul.cu
./matrixMul }
```

The Job Name

Your class account

This is the partition name that we use for classes

Double dash

bash code

Note: A single dash (-) is used for a single character flag
A double dash (--) is used for a word (more than a single character) flag

submit-c 143% sbatch submit.bash
Submitted batch job 474

Submit the job described in your shell file

submit-c 144% cat matrixmul.err

Check the output
(I like sending my output to standard error, not standard output)

Oregon State University
Computer Graphics

mjb - April 18, 2020

What Showed up in my Email

7

From	Subject
Slurm workload manager	Slurm Job_id=3980 Name=MatrixMul Ended, Run time 00:00:12, COMPLETED, ExitCode 0 2
Slurm workload manager	Slurm Job_id=3980 Name=MatrixMul Began, Queued time 00:00:01 1



mjb - April 18, 2020

Submitting a Loop

8

submitloop.bash:

```
#!/bin/bash
#SBATCH -J MatrixMul
#SBATCH -A cs475-575
#SBATCH -p class
#SBATCH --gres=gpu:1
#SBATCH -o matrixmul.out
#SBATCH -e matrixmul.err
#SBATCH --mail-type=BEGIN,END,FAIL
#SBATCH --mail-user=joeparallel@oregonstate.edu
for t in 1 2 4 8 16 32
do
    /usr/local/apps/cuda/cuda-10.1/bin/nvcc -DNUMT=$t -o matrixMul matrixMul.cu
    ./matrixMul
done
```

bash code

submit-c 153% sbatch submitloop.bash
Submitted batch job 475

submit-c 154% tail -f matrixmul.err

Displays the latest output added to matrixmul.err.
Keeps doing it forever.

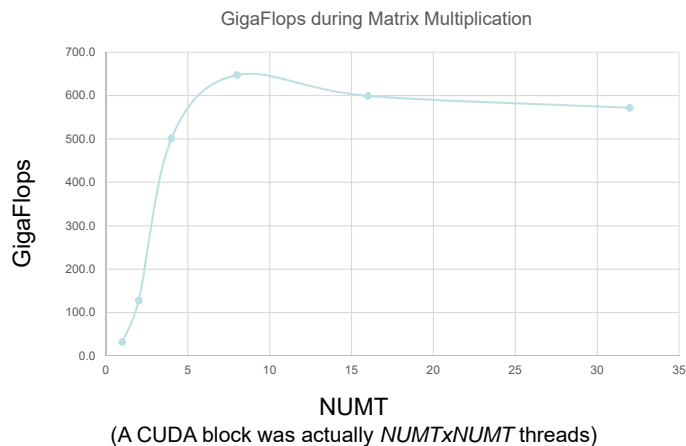
Control-c to get out of it.



mjb - April 18, 2020

Results for Multiplying two 1024x1024 Matrices

9



mjb - April 18, 2020

Use slurm's *scancel* if your Job Needs to Be Killed

10

submit-c 163% sbatch submitloop.bash
Submitted batch job 476

submit-c 164% scancel 476



mjb - April 18, 2020

Submitting an OpenCL job to the DGX Systems using Slurm

11

submit.bash:

```
#!/bin/bash
#SBATCH -J MatrixMult
#SBATCH -A cs475-575
#SBATCH -p class
#SBATCH --gres=gpu:1
#SBATCH -o printinfo.out
#SBATCH -e printinfo.err
#SBATCH --mail-type=BEGIN,END,FAIL
#SBATCH --mail-user=joeparallel@oregonstate.edu
{g++ -o printinfo printinfo.cpp /usr/local/apps/cuda/cuda-10.1/lib64/libOpenCL.so.1.1 -lm -fopenmp }
./printinfo
```



mjb - April 18, 2020


Here's What *printinfo* Got on the DGX System

12

```
Number of Platforms = 1
Platform #0:
  Name   = 'NVIDIA CUDA'
  Vendor = 'NVIDIA Corporation'
  Version = 'OpenCL 1.2 CUDA 10.1.351'
  Profile = 'FULL_PROFILE'
  Number of Devices = 1
Device #0:
  Type = 0x0004 = CL_DEVICE_TYPE_GPU
  Device Vendor ID = 0x10de (NVIDIA)
  Device Maximum Compute Units = 80
  Device Maximum Work Item Dimensions = 3
  Device Maximum Work Item Sizes = 1024 x 1024 x 64
  Device Maximum Work Group Size = 1024
  Device Maximum Clock Frequency = 1530 MHz
```

Device Extensions:

```
cl_khr_global_int32_base_atomics
cl_khr_global_int32_extended_atomics
cl_khr_local_int32_base_atomics
cl_khr_local_int32_extended_atomics
cl_khr_fp64
cl_khr_byte_addressable_store
cl_khr_icd
cl_khr_gl_sharing
cl_nv_compiler_options
cl_nv_device_attribute_query
cl_nv_pragma_unroll
cl_nv_copy_opts
cl_nv_create_buffer
```



mjb - April 18, 2020