

# CUDA Matrix Multiplication



**Oregon State  
University**

**Mike Bailey**

mjb@cs.oregonstate.edu



This work is licensed under a [Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License](https://creativecommons.org/licenses/by-nc-nd/4.0/)



**Oregon State  
University**  
Computer Graphics

cudaMatrixMult.pptx

mjb - March 9, 2020

## Anatomy of the CUDA *matrixMult* Program: #defines, #includes, and Globals

```
#include <stdio.h>
#include <assert.h>
#include <malloc.h>
#include <math.h>
#include <stdlib.h>

#include <cuda_runtime.h>
#include "helper_functions.h"
#include "helper_cuda.h"

#ifndef MATRIX_SIZE
#define MATRIX_SIZE 1024
#endif

#define AROWS      MATRIX_SIZE
#define ACOLS      MATRIX_SIZE

#define BROWS      MATRIX_SIZE
#define BCOLS      MATRIX_SIZE
#define ACOLSBROWS ACOLS      // better be the same!
#define CROWS      AROWS
#define CCOLS      BCOLS

float hA[AROWS][ACOLS];
float hB[BROWS][BCOLS];
float hC[CROWS][CCOLS];
```



**Oregon State  
University**  
Computer Graphics

mjb - March 9, 2020

### Anatomy of a CUDA Program: Error-Checking

3

```
void
CudaCheckError( )
{
    cudaError_t e = cudaGetLastError( );
    if( e != cudaSuccess )
    {
        fprintf( stderr, "CUDA failure %s:%d: '%s'\n", __FILE__, __LINE__, cudaGetErrorString(e));
    }
}
```



mjb - March 9, 2020

### Anatomy of a CUDA Program: The Kernel Function

4

```
__global__ void MatrixMul( float *A, float *B, float *C )
{
    // [A] is AROWS x ACOLS
    // [B] is BROWS x BCOLS
    // [C] is CROWS x CCOLS = AROWS x BCOLS

    int blockDim = blockDim.x * blockDim.y;
    int blockNum = blockIdx.x * blockDim.x + blockIdx.y * blockDim.y;
    int gid = blockNum + threadIdx.x * blockDim.x + threadIdx.y;

    int crow = gid / CCOLS;
    int ccol = gid % CCOLS;

    int aindex = crow * ACOLS; // a[i][0]
    int bindex = ccol; // b[0][j]
    int cindex = crow * CCOLS + ccol; // c[i][j]

    float cij = 0.;
    for( int k = 0; k < ACOLSBROWS; k++ )
    {
        cij += A[aindex] * B[bindex];
        aindex++;
        bindex += BCOLS;
    }
    C[cindex] = cij;
    __syncthreads( );
}
```



mjb - March 9, 2020

## Anatomy of a CUDA Program: Setting Up the Memory for the Matrices

5

// allocate device memory:

```
float *dA, *dB, *dC;  
cudaMalloc( (void **)&dA, sizeof(hA) );  
cudaMalloc( (void **)&dB, sizeof(hB) );  
cudaMalloc( (void **)&dC, sizeof(hC) );  
CudaCheckError( );
```

// copy host memory to device memory:

```
cudaMemcpy( dA, hA, sizeof(hA), cudaMemcpyHostToDevice );  
cudaMemcpy( dB, hB, sizeof(hB), cudaMemcpyHostToDevice );
```

This is a defined constant in one of the CUDA .h files

In `cudaMemcpy( )`, it's always the second argument getting copied to the first!



mjb - March 9, 2020

## Anatomy of a CUDA Program: Getting Ready to Execute

6

```
// setup execution parameters:  
dim3 threads( 16, 16, 1 );  
if( threads.x > CROWS )  
    threads.x = CROWS;  
if( threads.y > CCOLS )  
    threads.y = CCOLS;  
dim3 grid( CROWS / threads.x, CCOLS / threads.y );
```

```
// create cuda events for timing:  
cudaEvent_t start, stop;  
cudaEventCreate( &start );  
cudaEventCreate( &stop );  
CudaCheckError( );
```

```
// record the start event:  
cudaEventRecord( start, NULL );
```



mjb - March 9, 2020

## Anatomy of a CUDA Program: Executing the Kernel

7

```
// execute the kernel:
```

```
MatrixMul<<<grid, threads>>>( dA, dB, dC );
```

Function call arguments

# of blocks

# of threads per block

The call to **MatrixMul**( ) returns *immediately*!

If you upload the resulting array (dC) right away, it will have garbage in it.

To block until the kernel is finished, call:

```
cudaDeviceSynchronize( );
```



mjb - March 9, 2020

## Anatomy of a CUDA Program: Getting the Stop Time and Printing Performance

8

```
// record the stop event:
```

```
cudaEventRecord( stop, NULL );
```

```
// wait for the stop event to complete:
```

```
cudaEventSynchronize( stop );
```

```
float msecTotal;
```

```
cudaEventElapsedTime( &msecTotal, start, stop );
```

```
// performance:
```

```
float msecPerMatrixMul = msecTotal;
```

```
double flopsPerMatrixMul = (double)CROWS * (double)CCOLS * (double)ACOLSBROWS;
```

```
double gigaFlops = ( flopsPerMatrixMul / 1000000000. ) / ( msecPerMatrixMul / 1000.0 );
```

```
fprintf( stderr, "%6d\t%6d\t%10.3lf\n", CROWS, CCOLS, gigaFlops );
```



mjb - March 9, 2020

## Anatomy of a CUDA Program: Copying the Matrix from the Device to the Host

9

```
cudaMemcpy( hC, dC, sizeof(hC), cudaMemcpyDeviceToHost );  
CudaCheckError( );
```

```
// clean up:  
cudaFree( dA );  
cudaFree( dB );  
cudaFree( dC );  
CudaCheckError( );
```

This is a defined constant in one of the CUDA .h files

In **cudaMemcpy( )**, it's always the second argument getting copied to the first!