Performing Reductions in OpenCL

Mike Bailey

mjb@cs.oregonstate.edu

Computer Graphics

mjb – April 8, 2020

Recall the OpenCL Memory Model

Kernel

Global Memory

Constant Memory

WorkGroup

WorkGroup

WorkGroup

WorkGroup

Shared Memory

Work-Item

Work-Item

Work-Item

Private Memory

Private Memory

Private Memory
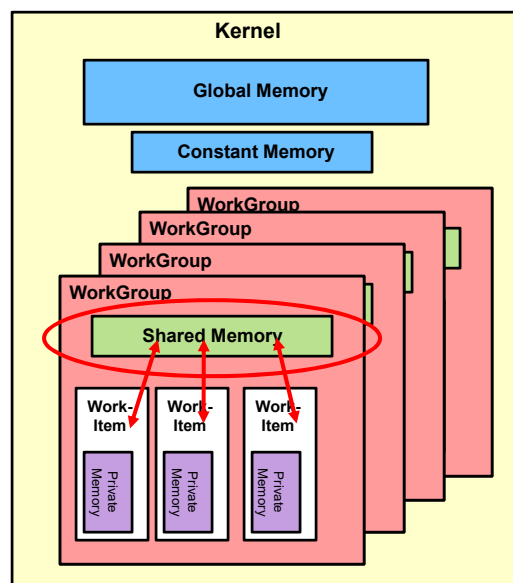
Computer Graphics

mjb – April 8, 2020

## Here's the Problem We are Trying to Solve

Like the *first.cpp* demo program, we are piecewise multiplying two arrays. Unlike the first demo program, we want to then add up all the products and return the sum.

$A * B \rightarrow prods$

$\Sigma\ prods \rightarrow C$

After the array multiplication, we want each work-group to sum the products within that work-group, then return them to the host in an array for final summing.

To do this, we will not put the products into a large global device array, but into a **prods[ ]** array that is shared within its work-group.
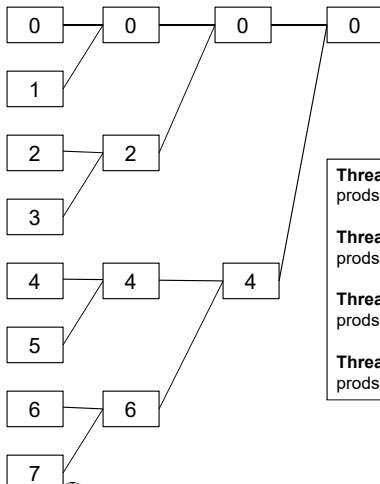
numItems = 8;

Computer Graphics

mjb – April 8, 2020

Reduction Takes Place in a Single Work-Group

numItems = 8;

If we had 8 work-items in a work-group, we would like the threads in each work-group to execute the following instructions . . .

Thread #0:	Thread #0:	Thread #0:
prods[0] += prods[1];	prods[0] += prods[2];	prods[0] += prods[4];
Thread #2:		
prods[2] += prods[3];		
Thread #4:	**Thread #4:**	
prods[4] += prods[5];	prods[4] += prods[6];	
Thread #6:		
prods[6] += prods[7];		

. . . but in a more general way than writing them all out by hand.

Computer Graphics

mjb – April 8, 2020

Here's What You Would Change in your Host Program

```
size_t  numWorkGroups = NUM_ELEMENTS  /  LOCAL_SIZE;

                • • •
float * hA  = new float [ NUM_ELEMENTS ];
float * hB  = new float [ NUM_ELEMENTS ];
float * hC  = new float [ numWorkGroups ];
size_t  abSize = NUM_ELEMENTS * sizeof(float);
size_t   cSize = numWorkGroups   * sizeof(float);

            • • •
cl_mem  dA = clCreateBuffer( context, CL_MEM_READ_ONLY,   abSize, NULL, &status );
cl_mem  dB = clCreateBuffer( context, CL_MEM_READ_ONLY,   abSize, NULL, &status );
cl_mem  dC = clCreateBuffer( context, CL_MEM_WRITE_ONLY,   cSize, NULL, &status );

            • • •
status = clEnqueueWriteBuffer( cmdQueue, dA, CL_FALSE, 0, abSize, hA, 0, NULL, NULL );
status = clEnqueueWriteBuffer( cmdQueue, dB, CL_FALSE, 0, abSize, hB, 0, NULL, NULL );

            • • •
cl_kernel kernel = clCreateKernel( program, "ArrayMultReduce", &status );
            • • •
status = clSetKernelArg( kernel, 0, sizeof(cl_mem),              &dA );
status = clSetKernelArg( kernel, 1, sizeof(cl_mem),              &dB );
status = clSetKernelArg( kernel, 2, LOCAL_SIZE * sizeof(float),  NULL );
                   // local "prods" array is dimensioned the size of each work-group
status = clSetKernelArg( kernel, 3, sizeof(cl_mem),              &dC );
```

$A * B \rightarrow$ prods

Σ prods $\rightarrow$ C

This NULL is how you tell OpenCL that this is a *local* (shared) array, not a global array

Computer Graphics

mjb – April 8, 2020

The Arguments to the Kernel

```
status = clSetKernelArg( kernel, 0, sizeof(cl_mem), &dA );
status = clSetKernelArg( kernel, 1, sizeof(cl_mem), &dB );
status = clSetKernelArg( kernel, 2, LOCAL_SIZE * sizeof(float), NULL );
                              // local "prods" array – one per work-item
status = clSetKernelArg( kernel, 3, sizeof(cl_mem), &dC );
```

```
kernel void
ArrayMultReduce( global const float *dA, global const float *dB, local float *prods, global float *dC )
{
        int gid        = get_global_id( 0 );     // 0 .. total_array_size-1
        int numItems = get_local_size( 0 );   // # work-items per work-group
        int tnum      = get_local_id( 0 );       // thread (i.e., work-item) number in this work-group
                                                 // 0 .. numItems-1
        int wgNum    = get_group_id( 0 );      // which work-group number this is in

        prods[ tnum ] =  dA[ gid ] * dB[ gid ];   // multiply the two arrays together

        // now add them up – come up with one sum per work-group
        // it is a big performance benefit to do it here while "prods" is still available – and is local
        // it would be a performance hit to pass "prods" back to the host then bring it back to the device for reduction
```

$A * B \rightarrow$ prods

Oregon State University
Computer Graphics

mjb – April 8, 2020

Reduction Takes Place Within a Single Work-Group
Each work-item is run by a single thread

Thread #0:
prods[0] += prods[1];

Thread #2:
prods[2] += prods[3];

Thread #4:
prods[4] += prods[5];

Thread #6:
prods[6] += prods[7];

offset = 1
mask = 1;

Thread #0:
prods[0] += prods[2];

Thread #4:
prods[4] += prods[6];

offset = 2
mask = 3;

Thread #0:
prods[0] += prods[4];

offset = 4
mask = 7;

A work-group consisting of *numItems* work-items can be reduced to a sum in Log_2*(numItems)* steps. In this example, *numItems=8*.

The reduction begins with the individual products in prods[0] .. prods[7].

The final sum will end up in prods[0], which will then be copied into dC[wgNum].

Oregon State
University
Computer Graphics

mjb – April 8, 2020

Reduction Takes Place in a Single Work-Group
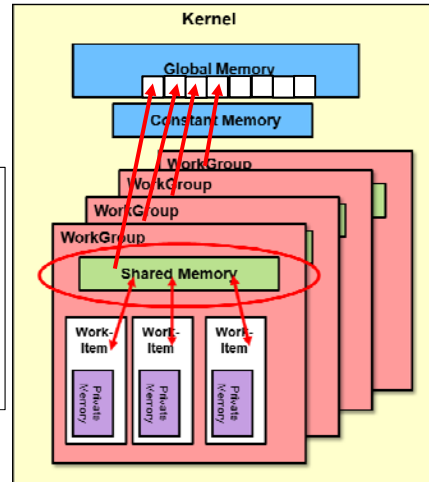### Each work-item is run by a single thread

Thread #0:
prods[0] += prods[1];

Thread #2:
prods[2] += prods[3];

Thread #4:
prods[4] += prods[5];

Thread #6:
prods[6] += prods[7];

offset = 1
mask = 1;

numItems = 8;

Thread #0:
prods[0] += prods[2];

Thread #4:
prods[4] += prods[6];

offset = 2
mask = 3;

Thread #0:
prods[0] += prods[4];

offset = 4
mask = 7;

```
kernel void
ArrayMultReduce( … )
    int gid       = get_global_id( 0 );
    int numItems = get_local_size( 0 );
    int tnum      = get_local_id( 0 );    // thread number
    int wgNum     = get_group_id( 0 );   // work-group number

    prods[ tnum ] =  dA[ gid ] * dB[ gid ];
```

```
// all threads execute this code simultaneously:
 for( int offset = 1; offset < numItems; offset *= 2 )
 {
     int mask = 2*offset - 1;
     barrier( CLK_LOCAL_MEM_FENCE );  // wait for completion
     if(  ( tnum & mask ) == 0 )
     {
         prods[ tnum ] += prods[ tnum + offset ];
     }
 }

barrier( CLK_LOCAL_MEM_FENCE );
if( tnum == 0 )
    dC[ wgNum ] = prods[ 0 ];
```

Σ prods $\rightarrow$ C

Oregon State
University
Computer Graphics

And, Finally, in your Host Program

```
Wait( cmdQueue );
double time0 = omp_get_wtime( );

status = clEnqueueNDRangeKernel( cmdQueue, kernel, 1, NULL, globalWorkSize, localWorkSize,
                                 0, NULL, NULL );
PrintCLError( status, "clEnqueueNDRangeKernel failed: " );

Wait( cmdQueue );
double time1 = omp_get_wtime( );


status = clEnqueueReadBuffer( cmdQueue, dC, CL_TRUE, 0, numWorkGroups*sizeof(float), hC,
                              0, NULL, NULL );
PrintCLError( status, "clEnqueueReadBufferl failed: " );
Wait( cmdQueue );


float sum = 0.;
for( int i = 0; i < numWorkgroups; i++ )
{
        sum += hC[ i ];
}
```

Oregon State
University
Computer Graphics

mjb – April 8, 2020

Reduction Performance
Work-Group Size = 32

Oregon State
University
Computer Graphics

mjb – April 8, 2020